

Football Corner Forecasting Report

Yujia Zhang, 21 May 2026

1 Summary

The baseline appears to be a Poisson-style model using the 10 anonymised match features plus the home indicator. This is a sensible starting point because corners are non-negative counts.

The final model keeps this baseline structure and adds one new signal: a pre-match estimate of each team’s attacking strength from historical corners. This signal is built with a Kalman filter, is strongly related to future corners, and the correlations with the anonymised features are low.

Rows with too little team history or no history at all need a fallback. Instead of filling the missing team-strength value with one global mean, a small regularised tree model predicts that value from the available match features. This gives a steady and reliable validation improvement.

More complex models and richer Kalman state splits were tested, but the dataset is too small for them to be reliable, likely leading to overfits. The simpler Poisson model plus one clear team-strength feature gave the best bias-variance tradeoff.

2 Final Model and Validation Results

The final model predicts the expected number of corners for each team in each match. It uses the 10 anonymised match features, the home indicator, and one extra team-strength feature: the team’s estimated attacking corner strength before the match.

The prediction model is a regularised Poisson GLM. Poisson is used because corners are non-negative counts. “Regularised” means the model is discouraged from using very large coefficients, which helps avoid overfitting.

$$\hat{y} \sim \text{PoissonRegressor}(\alpha = 0.1).$$

The value $\alpha = 0.1$ is the L2 penalty strength, selected on the tune seasons not on the validation seasons. Thus there’s no data leakage. small:

$$\text{Poisson count loss} + 0.1 \times \text{coefficient penalty}.$$

The best model is chosen as the one with the lowest loss:

$$0.25 \times \text{MAE} + 0.25 \times \text{RMSE} + 0.50 \times \text{Poisson deviance}.$$

The loss is the combination of three metrics that check different aspects of corner prediction quality. Poisson deviance receives the largest weight because corners are count data, so the model should generally fit the shape of a count outcome, not only minimise ordinary distance errors. MAE and RMSE are still included because the submitted output is a point prediction: MAE measures the typical miss in corners, while RMSE gives extra punishment to large misses.

The final table compares three models. The baseline is the provided reference prediction. The second row adds the Kalman attack-strength feature but fills missing values with a simple mean. The third row is the final model: it uses a tree-model fitted cold-start prior value instead of a mean fill when the sufficient team history is missing. Lower is better for all three metrics.

Model	MAE	RMSE	Poisson deviance
Baseline	2.2717	2.8574	1.3406
Poisson + mean-filled Kalman attack feature	2.1712	2.7349	1.2574
Poisson + cold-start Kalman attack prior	2.1588	2.7196	1.2437

3 Train-tune-validation Periods

All model choices are made before the final validation period:

$$\text{CORE} = 2012\text{-}2017, \quad \text{TUNE} = 2018\text{-}2019.$$

The final one-shot check is:

$$\text{TRAIN} = 2012\text{-}2019, \quad \text{VAL} = 2020\text{-}2022.$$

During validation feature construction, validation corners are masked before the Kalman filter is run. This means validation rows can receive pre-match states, but they cannot update any hidden state or home-advantage estimate.

4 Kalman Filter Team Strength

4.1 How the filter works

Each team has two hidden states:

$$\text{attack}_t = \text{corners the team tends to win,}$$

$$\text{defense}_t = \text{corners the team tends to allow the opponent to win.}$$




Figure 1: Kalman filter flow. The filter keeps one simple attacking signal for the final predictor; defence is used only to clean the update signal.

Raw corners are not a pure attack signal. A simple way to read a home team’s corner count is:

$$\text{home corners} \approx \text{home attack} + (\text{away defense} - \text{league mean}) + \text{home advantage.}$$

So before updating home attack, the filter subtracts the current opponent defense estimate and the league home advantage:

$$z_{\text{home attack}} = \text{home corners} - (\text{away defense} - \mu) - \text{home advantage.}$$

Here μ is the league-average corner level. If the away defense estimate equals the league mean it adds no adjustment; if it is above the mean that team allows extra opponent corners, and only that excess is removed from the raw home count. The same idea applies to away attack and to both defense states.

For each hidden state the Kalman update is:

$$K = \frac{P}{P + R}, \quad x_{\text{new}} = x_{\text{old}} + K(z - x_{\text{old}}), \quad P_{\text{new}} = (1 - K)P.$$

If the state is uncertain, the new match is trusted more; if the match observation is noisy, the update is more conservative.

Matches on the same date are batched: all pre-match features for that date are recorded first, then all hidden states are updated together. This avoids same-day leakage where an early kick-off could otherwise affect a later kick-off's features.

4.2 Why not a richer state space

The retained design keeps one attacking state and one defensive state per team, with home advantage handled as a league-level constant. A richer design was also tested, splitting each state into separate home and away versions: home attack, away attack, home defense, away defense. Conceptually this is attractive because teams can behave differently at home and away.




Figure 2: Richer Kalman design tested but not retained. More expressive, but it spreads the limited data across too many hidden states.

The cost is that each team now needs four hidden states instead of two, so each state receives fewer effective observations. With only 40 teams and 11 seasons the richer version had higher variance: it explained past data in more detail, but those details were not stable enough to improve future predictions. The simpler design is a better bias-variance tradeoff for this dataset.

4.3 Why not Elo-style ratings

An Elo-style model was also considered. Elo works well when one scalar rating can be updated from one match result. Corner prediction needs more structure: attacking strength, defensive weakness, home effect, and away effect. A multi-rating Elo would therefore require separate home attack, away attack, home defense, and away defense parameters per team, plus manual choices for the initial rating, update rates, home adjustment, and mean-reversion speed. With this dataset there is not enough validation data to tune all those choices reliably.

The Kalman filter handles updates more automatically. It tracks both an estimate and an uncertainty for each hidden state, so uncertain teams update more and stable teams update less. It requires only two hyperparameters (process noise and observation noise), both of which were selected on the tune period with a simple grid search.

4.4 Which Kalman signal to use

The filter exposes four pre-match signals per row: own attacking strength, own defensive weakness, opponent attacking strength, and opponent defensive weakness. Only the team's own attacking

strength is used in the final model. This was not chosen by intuition alone; it was verified empirically.

Kalman signal	Correlation with corners
Own attacking strength	0.2731
Own defensive weakness	-0.0145
Opponent attacking strength	-0.0152
Opponent defensive weakness	-0.0129

Table 1: Correlation on training rows after the Kalman warm-up period. Only the attacking-strength signal has a clear direct relationship to the target.

The attacking-strength signal is also not a duplicate of the 10 anonymised match features: its average absolute correlation with those features is 0.127 and the largest is 0.326, so it adds team-history information not already contained in the anonymised inputs.

Tune-period model	MAE	RMSE	Poisson deviance
No Kalman feature	2.1683	2.6954	1.3753
Add own attacking strength	2.1399	2.6496	1.3379
Add attacking and defensive strength	2.1404	2.6498	1.3380
Add all four Kalman signals	2.1465	2.6575	1.3439

Table 2: Tune-period ablation for Kalman feature selection. Adding the attacking-strength signal improves the model; adding the other signals does not improve it further.

One additional Kalman feature gives a clear tune-period gain, while adding the remaining signals introduces complexity without a robust improvement.

5 Cold-Start Prior

5.1 Why the problem is significant

The Kalman filter requires a warm-up period before its estimates are reliable. A team’s `kf.attack` value is only recorded once *both* that team and its opponent have each completed at least 5 matches. On the training set this affects 185 matches (370 rows, 8.9 % of training data) and is tolerable. On the holdout set the problem is more serious.

Four of the 28 holdout teams—`team_30`, `team_34`, `team_36`, and `team_47`—never appear in training at all, so they carry zero Kalman history into the holdout period. Because the warm-up condition requires *both* teams in a match to have enough history, every match involving a new team produces two missing rows: one for the new team and one for its opponent. In total, **360 of the 1,520 holdout rows (23.7 %) have no `kf.attack` value.**

5.2 Why a global mean is not enough

A global-mean fill would assign every cold-start team the same average attack strength, ignoring all 10 anonymised match features and the home indicator. For a completely new team, those features are the only signal about how corner-prone it is likely to be. Different teams have different feature

profiles and plausibly different attacking strengths; collapsing them to one number discards that signal.

5.3 What the prior model does

The prior is a supervised model for the *missing* Kalman attack feature, not a corner-count model. It is fitted only on training rows where the Kalman filter has already produced a real pre-match attacking-strength estimate, and it learns the mapping

$$(\text{feature_1}, \dots, \text{feature_10}, \text{home}) \longrightarrow \text{kf_attack}.$$

At inference time it fills only the rows where the Kalman attack estimate is missing. Rows with genuine Kalman history keep their original value untouched; the prior is purely a fallback for cold-start rows. The final Poisson model receives the same set of features for every row: the 10 anonymised match features, the home indicator, and the attacking-strength estimate—drawn from real Kalman history where available, and from the prior otherwise. This preserves the historical Kalman signal wherever it exists, while giving cold-start rows a feature-conditioned initial estimate rather than a single league-average number.

5.4 Conservative clipping

After the prior produces an initial attack value, that value is clipped to the 5th–95th percentile of the observed `kf_attack` distribution from training rows with real Kalman history. This prevents the prior from assigning implausibly large or small values for feature combinations it has not seen. The prior can still distinguish stronger from weaker cold-start rows within that range; it simply cannot extrapolate to an extreme initial strength unsupported by any historical Kalman state.

5.5 Model choice and hyperparameters

Ridge regression and `HistGradientBoostingRegressor` were both tested as the prior model on the tune period. Ridge assumes a linear relationship between the anonymised features and `kf_attack`; it also requires the features to be on comparable scales, which cannot be verified for anonymous inputs. `HistGradientBoostingRegressor` (HGB) makes neither assumption: it can capture non-linear feature-to-attack relationships and handles mixed-scale inputs natively without preprocessing.

Prior model	MAE	RMSE	Poisson deviance
Ridge ($\alpha = 1$)	2.1522	2.6765	1.3552
Ridge ($\alpha = 10$)	2.1533	2.6779	1.3564
Ridge ($\alpha = 100$)	2.1583	2.6849	1.3623
HGB (selected)	2.1460	2.6685	1.3503

Table 3: Tune-period comparison of cold-start prior models, evaluated on all tune rows including cold-start rows. HGB outperforms the best Ridge variant on every metric.

HGB outperforms the best Ridge variant on every metric. `HistGradientBoostingRegressor` is therefore used because it handles mixed-scale anonymous features without manual preprocessing and can capture non-linear relationships between the features and attack strength.

The parameters are set to be deliberately conservative rather than aggressively optimised:

Parameter	Value	Reason
<code>max_iter</code>	100	Limits ensemble size
<code>max_leaf_nodes</code>	15	Keeps individual trees small and shallow
<code>min_samples_leaf</code>	30	Roughly 1 % of training rows per leaf
<code>l2_regularization</code>	5.0	Strong weight shrinkage
<code>learning_rate</code>	0.03	Slow, stable updates

Table 4: Cold-start prior hyperparameters. All settings favour stability over fitting power, because the prior is a fallback rather than the main model.

The hyperparameters were set by hand to produce stable, non-extreme estimates. Aggressive tuning would risk overfitting the features-to-attack mapping on a relatively small training set, which would defeat the purpose of being conservative. The benefit of the prior over a plain mean fill was confirmed on the tune period before the validation period was touched: mean-fill gave a tune MAE of 2.1638, while the prior-fill gave 2.1460. The same ordering held on the one-shot validation check (MAE 2.1712 versus 2.1588).

5.6 Pipeline integration

In the final prediction pipeline the Kalman filter is first run over all labelled rows concatenated with the holdout rows. Labelled rows update the hidden states; holdout rows receive pre-match states but do not update anything because their corners are unknown. The Poisson corner model is fitted only on labelled rows that have real Kalman history. The cold-start prior is then fitted on those same rows and applied to fill missing `kf_attack` values in the holdout. The prior is therefore used only as a targeted fallback for missing team-history features, not as a replacement for either the Kalman filter or the final corner model.

6 Removed Experiments

Each experiment below was evaluated on the tune period under the same criteria. None improved the result robustly; the reasons fall into three categories.

6.1 Redundant signal: already captured by existing features

Static team target encoding. Each team’s historical mean corners scored and conceded were added as extra features. These are cruder versions of `kf_attack`: the Kalman filter already summarises each team’s attacking history, updating that estimate match by match with uncertainty tracking. The static mean ignores recency and carries no uncertainty signal. In ablation it added nothing after `kf_attack` was already in the model.

Rolling form features. Short-window rolling means (last 3, 5, and 10 matches) overlap substantially with what `kf_attack` already captures. The Kalman filter is itself a weighted average of past observations, emphasising recent matches through its uncertainty dynamics. Adding rolling means alongside it therefore introduces near-duplicate columns. The tune-period ablation showed no consistent gain from any rolling window, alone or combined with `kf_attack`.

Team, opponent, and season identity features. Team and opponent identities are already encoded through `kf_attack`. Adding them as categorical dummy variables asks the model to re-learn a static per-team offset that the Kalman feature already provides dynamically. Season identity produced a scale problem: the raw year values (2012–2022) are an order of magnitude larger than the other features, which destabilised the Poisson GLM and caused convergence issues even after regularisation.

Frozen home/away drift features. Separate home-only and away-only rolling averages were considered to capture venue-specific form. After controlling for the `home` indicator and `kf_attack`, the residual home/away drift signal was too small to recover reliably in the tune period.

6.2 Too many parameters for the dataset size

The training set contains around 4,180 rows (2,090 matches $\times$ 2 rows). Any model that introduces many new parameters to estimate risks high variance: it can fit the training data well but generalises poorly.

Rich Kalman filter with separate home and away states. The simpler Kalman design keeps two states per team (attack and defence). A richer version was also tested that splits each into a home and an away state, giving four states per team. With 40 teams this means roughly 160 latent variables to track instead of 80. Each state receives fewer effective observations, so the estimates are noisier. The richer design produced a worse honest validation result (VAL MAE 2.276 versus 2.159) despite being more expressive in principle.

Tree models, random forest, GBM, and XGBoost as the main corner model. Non-linear tree ensembles were tested with a Poisson or Tweedie objective. With only around 4,000 training rows, these models need strong regularisation and careful cross-validation to avoid overfitting. Even with early stopping and depth limits, none improved reliably over the simple Poisson GLM on the tune period. The anonymised features also provide no natural interaction structure to exploit, so the non-linear capacity was not needed.

Semi-linear and GAM-style Poisson models. Generalised additive models allow some features to enter through smooth spline terms rather than linearly. Each spline term adds effective degrees of freedom. With the current feature set and dataset size, this extra flexibility did not help: the spline fits on CORE generalised poorly to TUNE.

Tweedie and Negative Binomial GLMs. Corner counts show genuine overdispersion: the marginal variance-to-mean ratio is approximately 1.84, which in principle motivates a Tweedie or Negative Binomial distribution over plain Poisson. Despite this theoretical motivation, neither improved over Poisson on the tune period. Three factors explain why.

First, each model introduces one additional global parameter to estimate alongside the regularisation strength: the power $p \in (1, 2)$ for Tweedie, or the dispersion ϕ for Negative Binomial. Tuning two coupled parameters reliably requires more data than the roughly 4,000 training rows provide.

Second, the marginal variance-to-mean ratio of 1.84 reflects between-team heterogeneity—strong attacking teams consistently win more corners. This source of overdispersion is already

handled by the team-strength feature from the Kalman filter. The residual overdispersion after conditioning on all features is likely considerably smaller than the marginal figure.

Third, the composite tuning metric weights Poisson deviance most heavily. Tweedie and Negative Binomial models are optimised under their own likelihoods, which are not the same criterion. Even a distributional improvement can go undetected if the chosen evaluation measure is calibrated to Poisson.

7 What Next

Five directions below

7.1 Non-linear team-strength response

The Poisson GLM assumes a log-linear relationship between the attacking-strength feature and expected corners. This may be too rigid: a team whose attack estimate is far above average might produce corners at a rate that grows faster than the linear term implies. A Poisson GAM with a smooth spline on the attacking-strength feature and linear terms for the base features would test whether a non-linear response adds predictive value.

7.2 Season-boundary state decay

The Kalman filter currently carries team states unchanged from the last match of one season to the first match of the next. Teams change substantially over a summer: new signings, departures, new manager, pre-season form. Shrinking each state towards the league mean at each season boundary—with a tunable decay parameter—would better reflect this turnover.

7.3 Opponent-aware cold-start prior

The current prior predicts a team’s attacking strength from its own match features only. When the prior is applied, the opponent’s Kalman attack and defence estimates are often already available. Adding those as inputs to the prior would make cold-start estimates sensitive to match context: a new team playing a historically strong defensive side should receive a lower initial attacking estimate than the same team playing a weak one.

7.4 Rolling cross-validation for hyperparameters

The process-noise and observation-noise grid search and the regularisation sweep both rely on a single CORE to TUNE split. Rolling-window cross-validation—for example training on 2012–2015 and predicting 2016–2017, then shifting the window forward by two years—would give more robust hyperparameter estimates and reduce the risk that the selected values happen to fit one particular two-year tune period well.